جامعة خليفة
Khalifa University

# Research Computing Introductory Session

Dr. Sergio Martinez

# What is the HPC cluster? - Some concepts

- **Node (Host):** Refers to the physical machine/server. In current systems, a node would typically include one or more processors, as well as memory and other hardware.

- **Processor:** Refers to the central processing unit (CPU), which contains one or more cores.

- **Core:** Refers to the basic computation unit of the CPU. This is unit that carries out the actual computations.

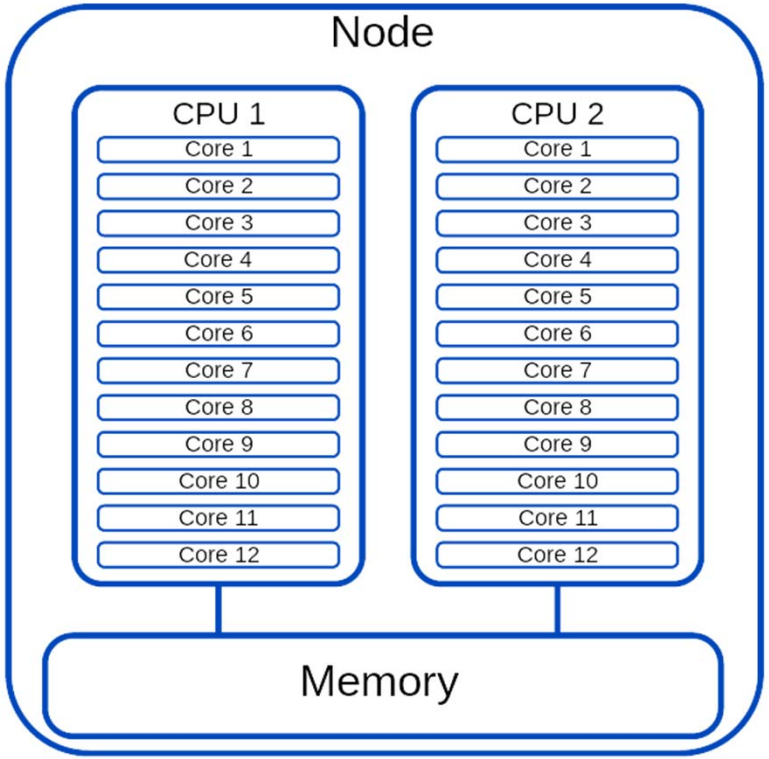

# What is the HPC cluster? – Almesbar specifications

- **Number of nodes**
- Cores and memory
- Network and storage
- Performance

| Number of nodes | |
|---|---|
| 88 CPU nodes | 2x Intel Xeon Gold 6230R 26-Core 2.1GHz (Cascade lake) (Q1‘20) |
| 3 GPU nodes | 2x Intel Xeon Gold 6230R 26-Core 2.1GHz(Cascade lake) (Q1’20) |
| | 4x NVidia V100 |

# What is the HPC cluster? – Almesbar specifications

- Number of nodes
- **Cores and memory**
- Network and storage
- Performance

| Cores per node | |
|---|---|
| 52 | |
| **Memory per node** | |
| CPU nodes | 384 GB |
| GPU nodes | 640 GB + 128 GB (GPU) |

| Total CPU cores |
|---|
| 4,732 |
| **Total GPU cores** |
| 61,440 |
| **Total memory** |
| 35.25 TB |

# What is the HPC cluster? – Almesbar specifications

- Number of nodes
- Cores and memory
- **Network and storage**
- Performance

| Network interconnection | |
|---|---|
| Infiniband EDR (100Gb/s) and Infiniband HDR (200Gb/s) | |
| **Storage** | |
| 300 TB home | NFSv4 (10Gb/s) |
| 947 TB scratch | Parallel file system (Lustre) (200Gb/s) |

# What is the HPC cluster? – Almesbar specifications

- Number of nodes
- Cores and memory
- Network and storage
- **Performance**

| Theoretical peak of performance (CPU) |
| --- |
| 307.5 TFOPS |
| **Achieved peak of performance (CPU)** |
| 205.0 TFOPS |

# What is the HPC cluster? - Layout

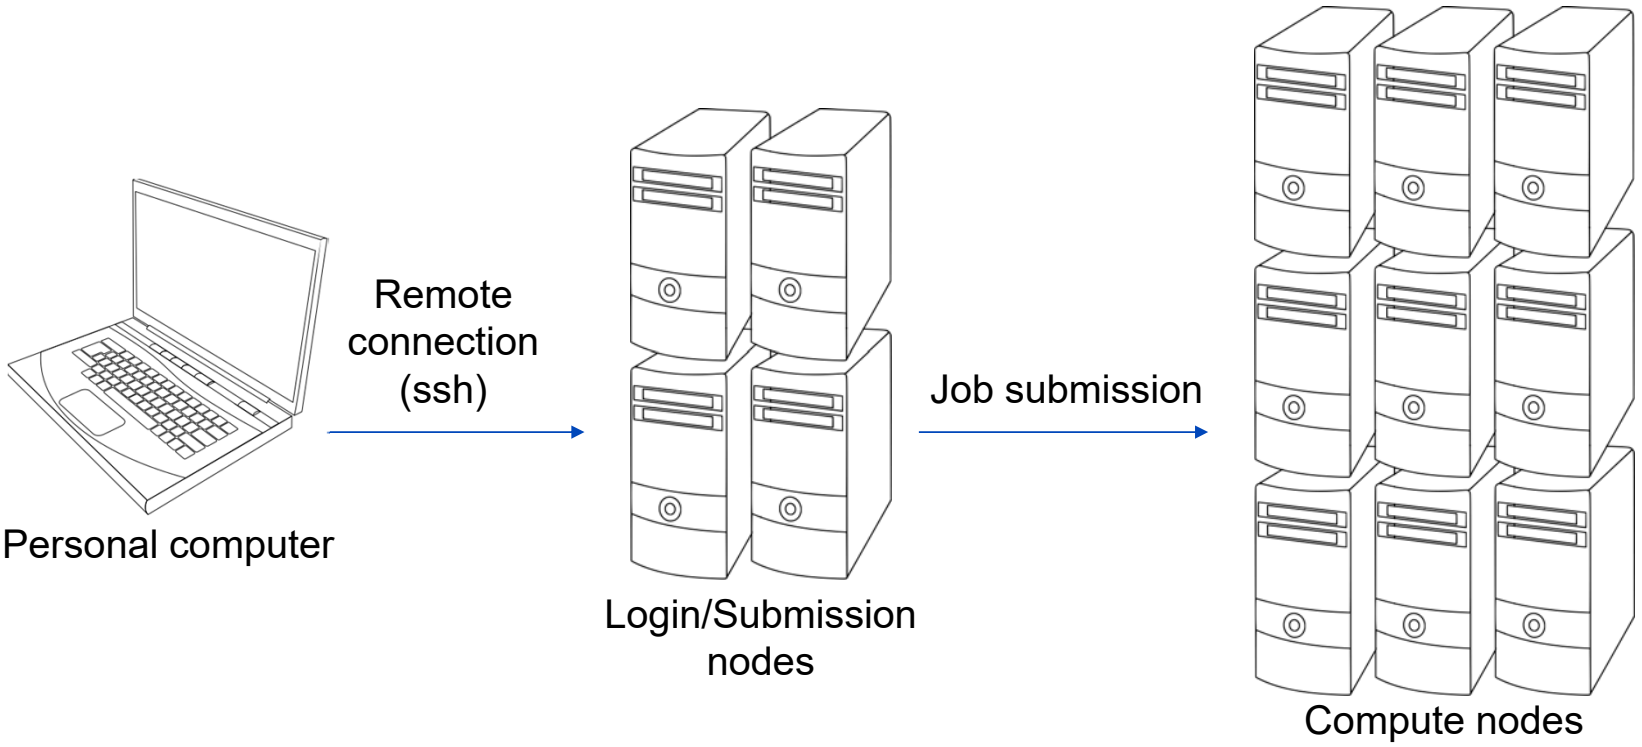

# What is the HPC cluster? - Personal computer

- Develop testing code (IDE). ✓
- Prepare simulations. ✓
- Remote connection with the cluster. ✓
- File transfer with the cluster. ✓

# What is the HPC cluster? - Submission nodes

- Organize files (create, delete, move, copy).✓
- Create scripts. ✓
- Compile small programs.✓
- Run applications.✗
- Connect with compute nodes.✗
- Submit jobs.✓

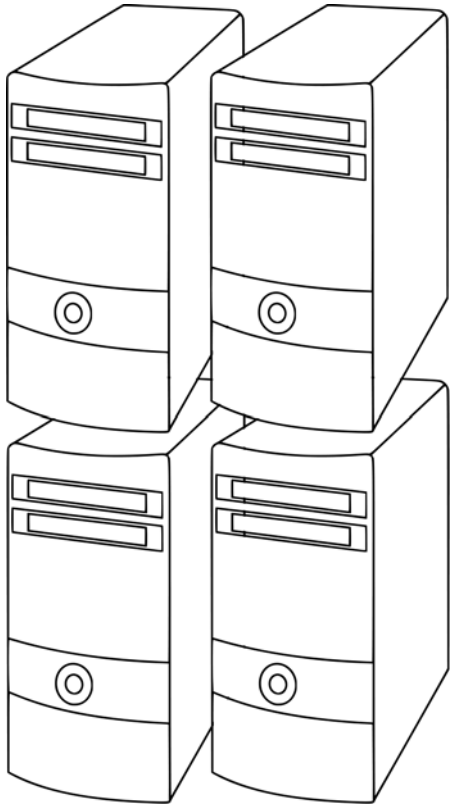

# What is the HPC cluster? - Compute nodes

- Run applications directly. ✗
- Connect from other node. ✗
- Fully controlled by the scheduler ✓

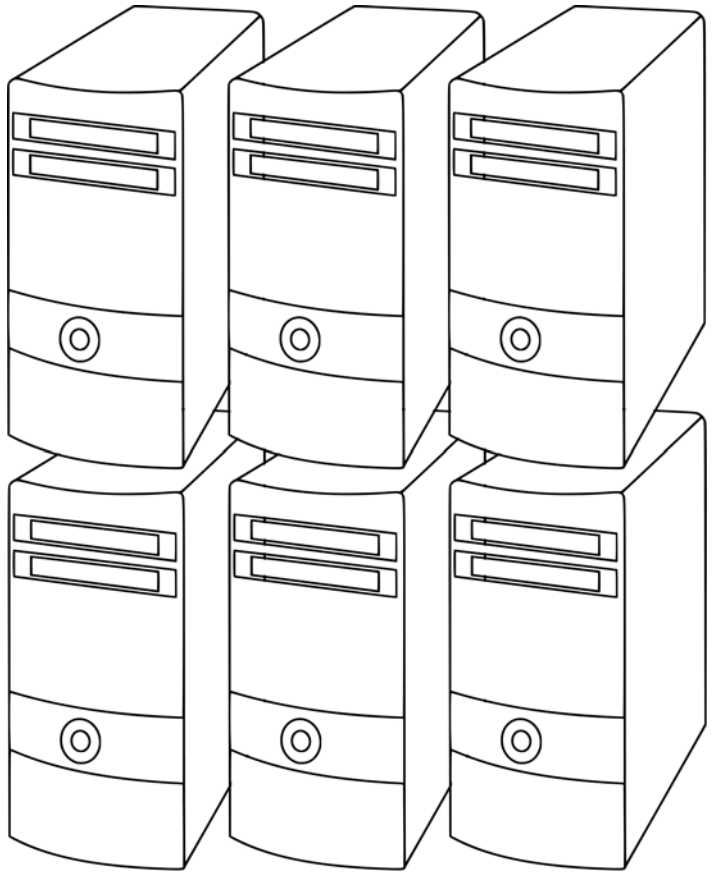

# What is the HPC cluster? - Shared facility

- Although the number of available resources on the cluster is quite high, the number of active users is also high.
- There is no limitation on the number of resources that one user can request. We cannot guarantee that the requested resources will be available at the moment of submitting a job.
- Jobs will remain pending in a queue until all the requested resources are available.
- Those resources are managed by the scheduler.
- Once the resources are available, jobs will start running on those resources.

# What is the HPC cluster? - Jobs

- Job is the unit of execution on the HPC cluster.
- Instead of running applications directly on a submission or on a compute node, in order to make an efficient use of the resources, we encapsulate our application in a job and we submit it through the scheduler.
- The syntax of a job allows us, among other things, to specify with accuracy what are the resources needed for our program to run and for how long we need those resources.

```
run python script
```

```
Job start
        Cores = 2
        Nodes = 1
        Queue = gpu
        Walltime = 03:00
        run python script
Job end
```

# What is the HPC cluster? - Scheduler

- The scheduler is a program that understands the job syntax and allocates the resources accordingly.
- The scheduler is configured in a way that enhances both the system and the users performance.
- It can manage multiple queues where jobs are submitted to.

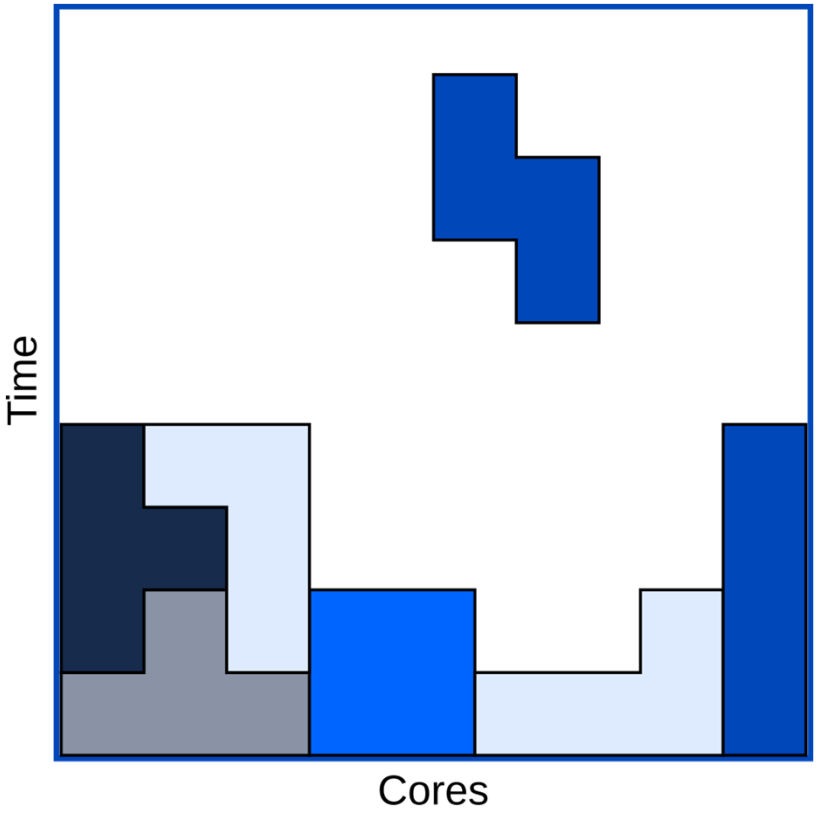

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- **Batch and Interactive jobs.**
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- Backfill
- Fairshare (dynamic priority)

- **Batch jobs:** run in the background. The don't require user interaction.
- **Interactive jobs:** allow input and output using command line interface or graphical user interface.

Interactive jobs are not recommended on a HPC cluster for several reasons:

- User might have to wait several hours before the application's GUI opens.
- Graphics need to travel over the network, causing a bad user experience (clunky interface)

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- **Array jobs**
- Exclusive mode
- Default wall time
- Maximum wall time
- Backfill
- Fairshare (dynamic priority)

- **Array syntax:** is an efficient way of submitting multiple jobs simultaneously
- Avoids “looping” and unnecessarily submitting a large amount of jobs with different JOBIDs.
- Ensemble solutions
- Parameter sweep

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- **Exclusive mode**
- Default wall time
- Maximum wall time
- Backfill
- Fairshare (dynamic priority)

- **Exclusive mode:** guarantees that the nodes allocated for the jobs to run will not be shared by other jobs.
- In other words, the scheduler cannot place an exclusive job unless there is a host that has no jobs running on it.
- Useful for memory consuming applications, regardless the number of cores needed.

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- **Default wall time**
- Maximum wall time
- Backfill
- Fairshare (dynamic priority)

In order to enhance the system performance, the scheduler needs to know the exact wall time of every single job in the system. This have a strong impact on dispatch planification, enabling features like backfill.

- Default wall time: if the user doesn't specify the wall time of a job at submission time, the scheduler will assume a default one for that job.
- The more accurate an user is specifying the wall time of their jobs, the lesser their jobs will be pending in the queue.

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- **Maximum wall time**
- Backfill
- Fairshare (dynamic priority)

In order to enhance the system performance, the scheduler needs to know the exact wall time of every single job in the system. This have a strong impact on dispatch planification, enabling features like backfill.

There is a limitation on the maximum wall time that a job can be running. If this limitation wasn't there:

- Overall pending time would increase notably.
- Non checkpointed jobs would be harmed if an error happens after running for long time.

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- **Maximum wall time**
- Backfill
- Fairshare (dynamic priority)

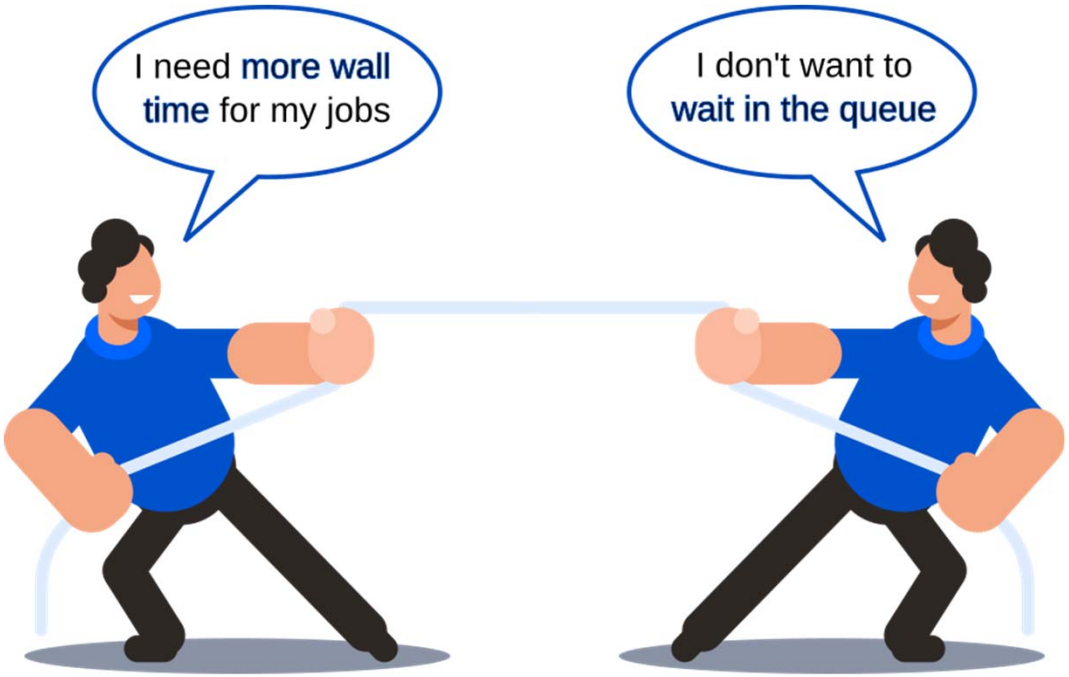

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- **Backfill**
- Fairshare (dynamic priority)

The scheduler knows the wall time for all the jobs submitted to a queue.

Large jobs that cannot start running because lack of resources are scheduled for a specific time in the future and they wait in the queue.

At the same time, as job finish their execution, resources become available. Those resources became an opportunity for small jobs to start running.

Backfill allows small jobs to jump in the queue and use those resources as long as they finish before large jobs start.

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- **Backfill**
- Fairshare (dynamic priority)

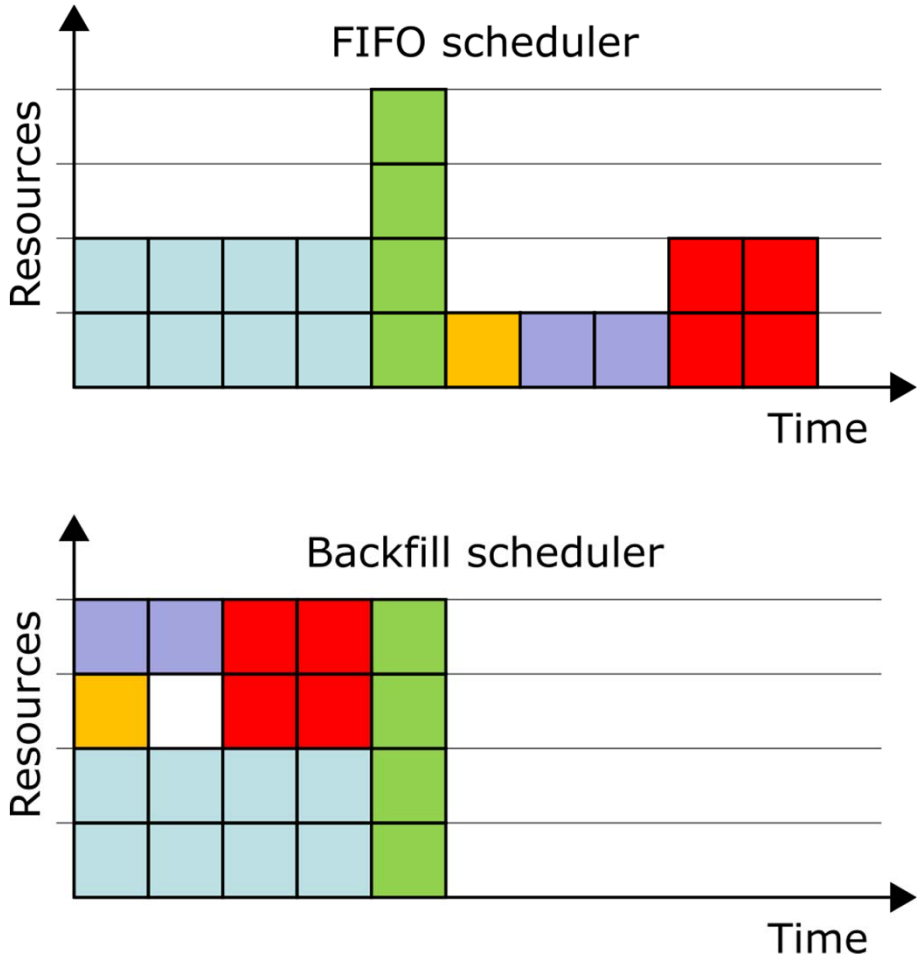

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- **Backfill**
- Fairshare (dynamic priority)

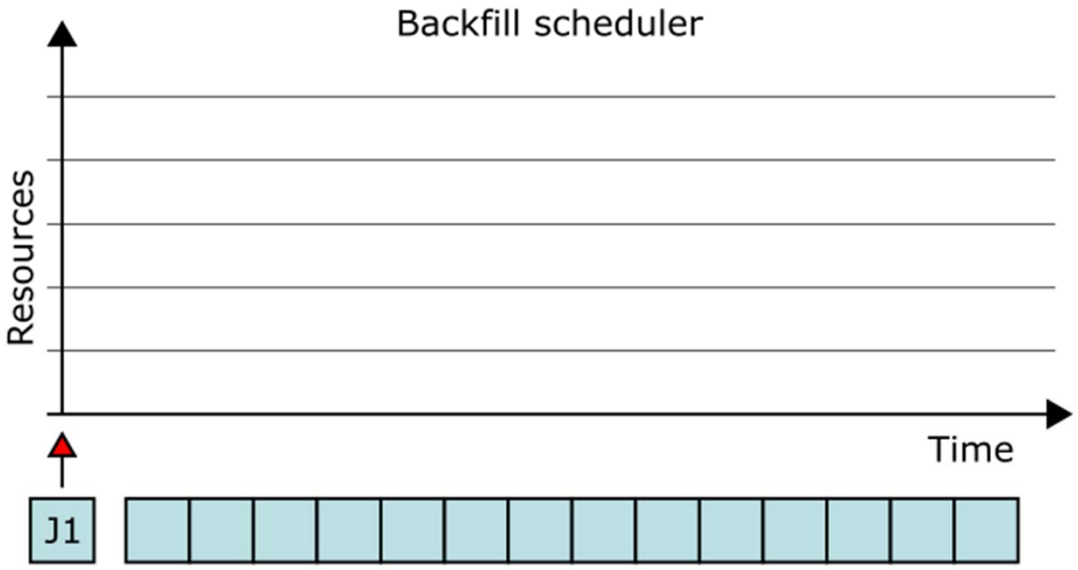

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- **Backfill**
- Fairshare (dynamic priority)

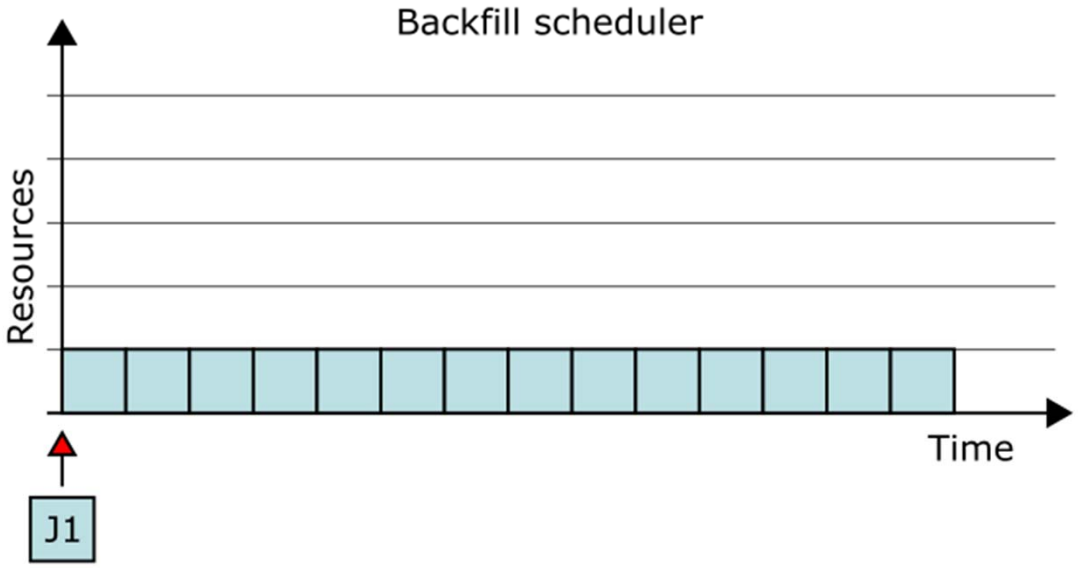

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- **Backfill**
- Fairshare (dynamic priority)

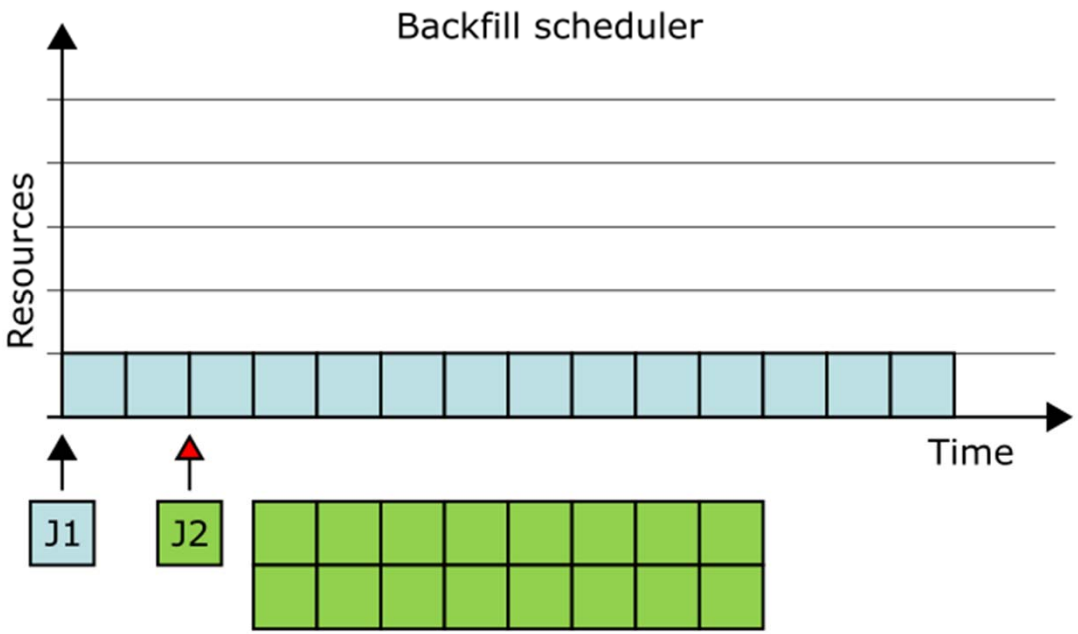

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- **Backfill**
- Fairshare (dynamic priority)

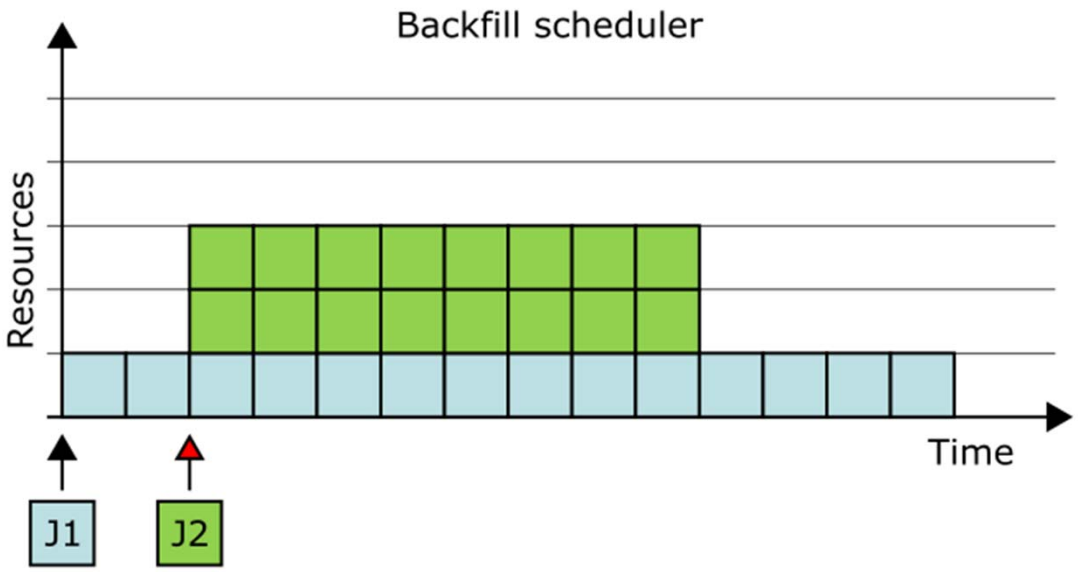

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- **Backfill**
- Fairshare (dynamic priority)

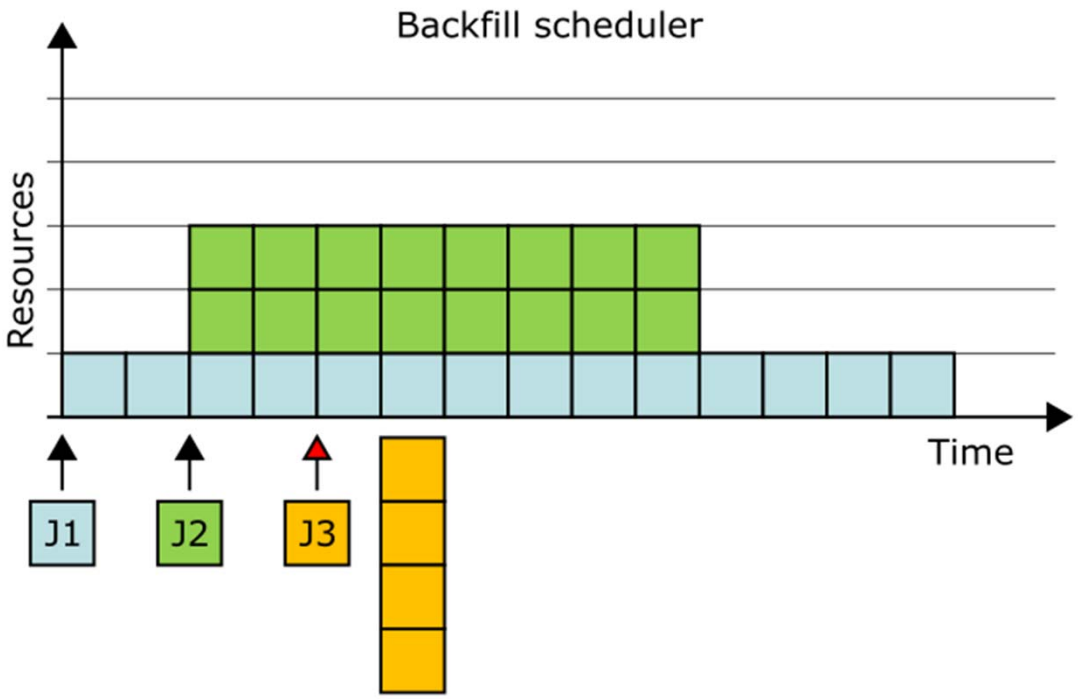

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- **Backfill**
- Fairshare (dynamic priority)

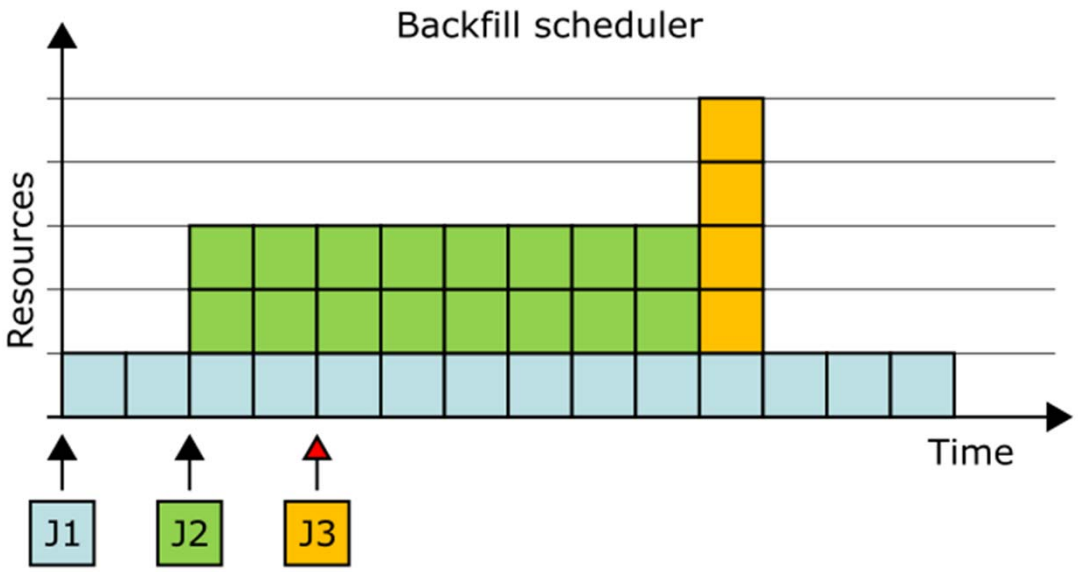

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- **Backfill**
- Fairshare (dynamic priority)

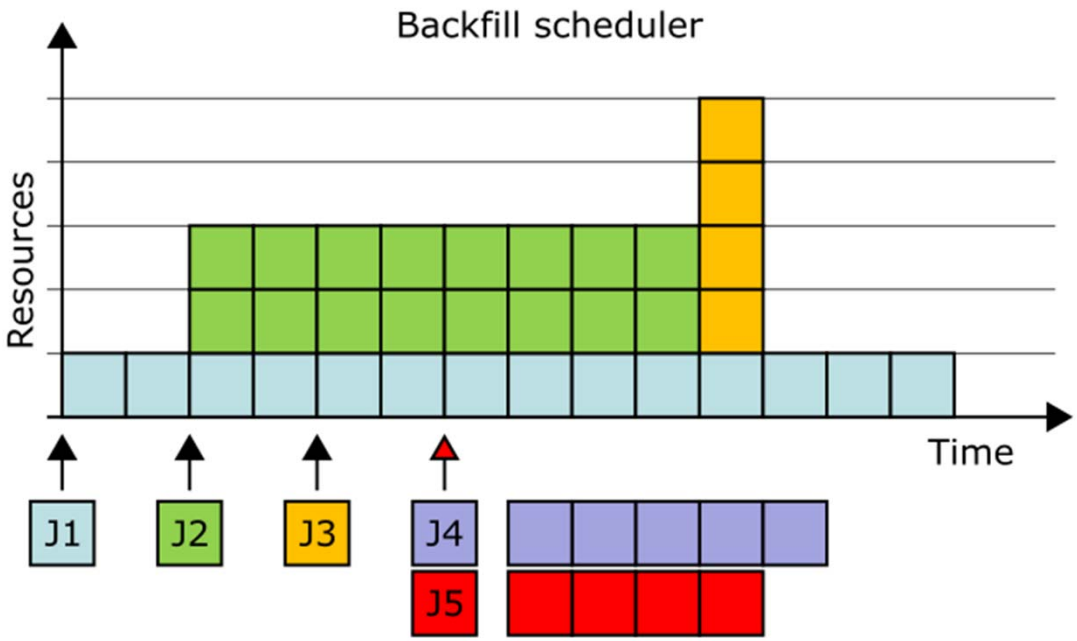

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- **Backfill**
- Fairshare (dynamic priority)

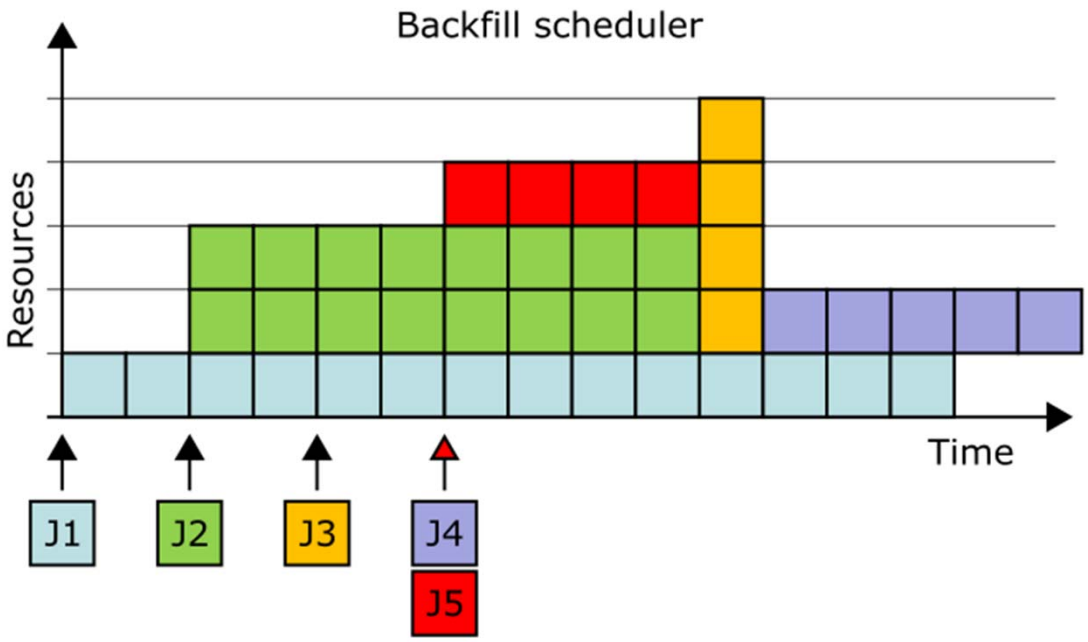

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- Backfill
- **Fairshare (dynamic priority)**

In order to enhance the user performance, fairshare feature calculates and assigns users' or projects' priority. This priority is inversely proportional to the utilization of the resources.

It is calculated first at project level and then at user level within the same project.

If the amount of resources used by all the jobs associated with Project 1 is higher than the amount of resources used by all the jobs associated with Project 2, then jobs submitted against Project 2 will have higher priority than jobs submitted against Project 1 and, hence, they will be dispatched sooner.

At the same time within the same project, if User 1 uses more resources then User 2, then User 2 will have higher priority than User 1 and, hence, their jobs will be dispatched sooner.

# What is the HPC cluster? - Queues

Queues are affected by different features and limitations:

- Batch and Interactive jobs.
- Array jobs
- Exclusive mode
- Default wall time
- Maximum wall time
- Backfill
- **Fairshare (dynamic priority)**

# What can you expect from the HPC cluster?

- **Solve problems that don't fit in a single node memory space.**
- Solve problems that can't be solved in a reasonable time.
- Solve problems that require large computational power using multiple nodes.
- Solve large amount of problems at the same time.
- Large amount of applications and libraries centrally installed and maintained.

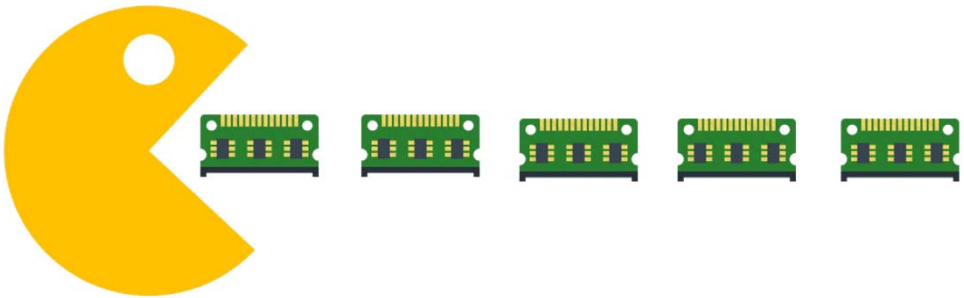

# What can you expect from the HPC cluster?

- Solve problems that don't fit in a single node memory space.
- **Solve problems that can't be solved in a reasonable time.**
- Solve problems that require large computational power using multiple nodes.
- Solve large amount of problems at the same time.
- Large amount of applications and libraries centrally installed and maintained.

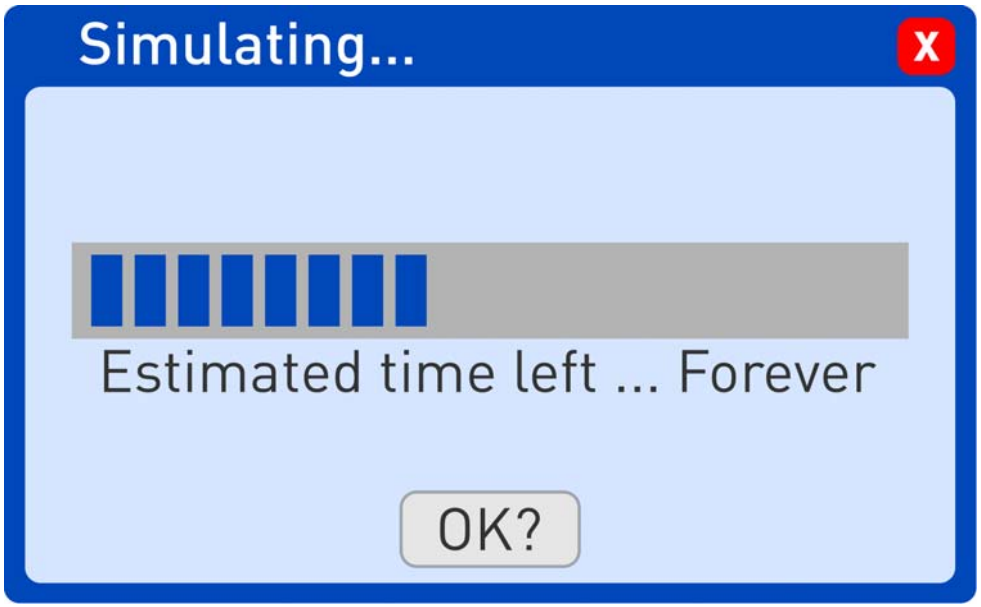

# What can you expect from the HPC cluster?

- Solve problems that don't fit in a single node memory space.
- Solve problems that can't be solved in a reasonable time.
- **Solve problems that require large computational power using multiple nodes.**
- Solve large amount of problems at the same time.
- Large amount of applications and libraries centrally installed and maintained.

# What can you expect from the HPC cluster?

- Solve problems that don't fit in a single node memory space.
- Solve problems that can't be solved in a reasonable time.
- Solve problems that require large computational power using multiple nodes.
- **Solve large amount of problems at the same time.**
- Large amount of applications and libraries centrally installed and maintained.

# What can you expect from the HPC cluster?

- Solve problems that don't fit in a single node memory space.
- Solve problems that can't be solved in a reasonable time.
- Solve problems that require large computational power using multiple nodes.
- Solve large amount of problems at the same time.
- **Large amount of applications and libraries centrally installed and maintained.**






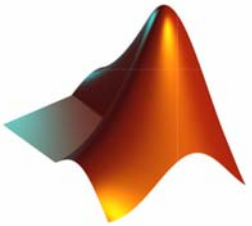

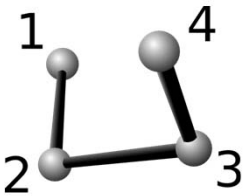



# What do you need to know in order to use it?

- **Basic Linux commands.**
- Scheduler.
- Theoretical limits.
- Practical limits.
- Responsible usage.

- Command line interface.
- Remote connection (ssh).
- Navigating files and directories (cd, ls).
- Working with files and directories (mkdir, cp, mv rm).
- Shell scripts.

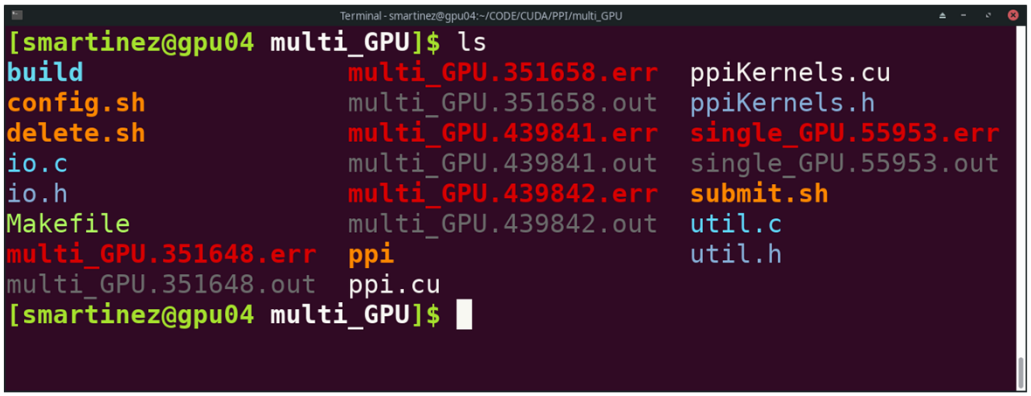

# What do you need to know in order to use it?

- Basic Linux commands.
- **Scheduler.**
- Theoretical limits.
- Practical limits.
- Responsible usage.

Each **HPC cluster** has its own **scheduler** and hence its own **syntax**:

| | **SLURM (Almesbar HPC cluster)** |
|---|---|
| **Submit job** | sbatch |
| **Display resources** | sinfo |
| **Monitor jobs** | scontrol |
| **Kill jobs** | scancel |

# What do you need to know in order to use it?

- Basic Linux commands.
- Scheduler.
- **Theoretical limits.**
- Practical limits.
- Responsible usage.

# What do you need to know in order to use it?

- Basic Linux commands.
- Scheduler.
- **Theoretical limits.**
- Practical limits.
- Responsible usage.

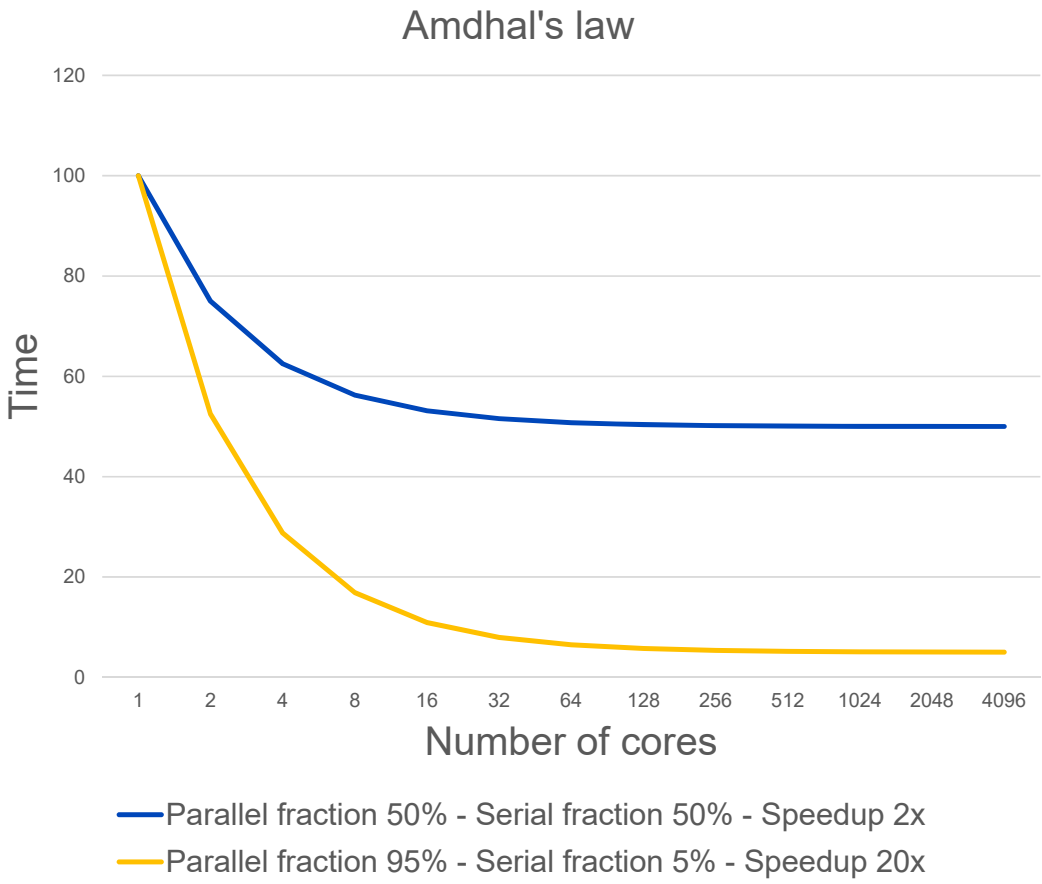

# What do you need to know in order to use it?

- Basic Linux commands.
- Scheduler.
- Theoretical limits.
- **Practical limits.**
- Responsible usage.

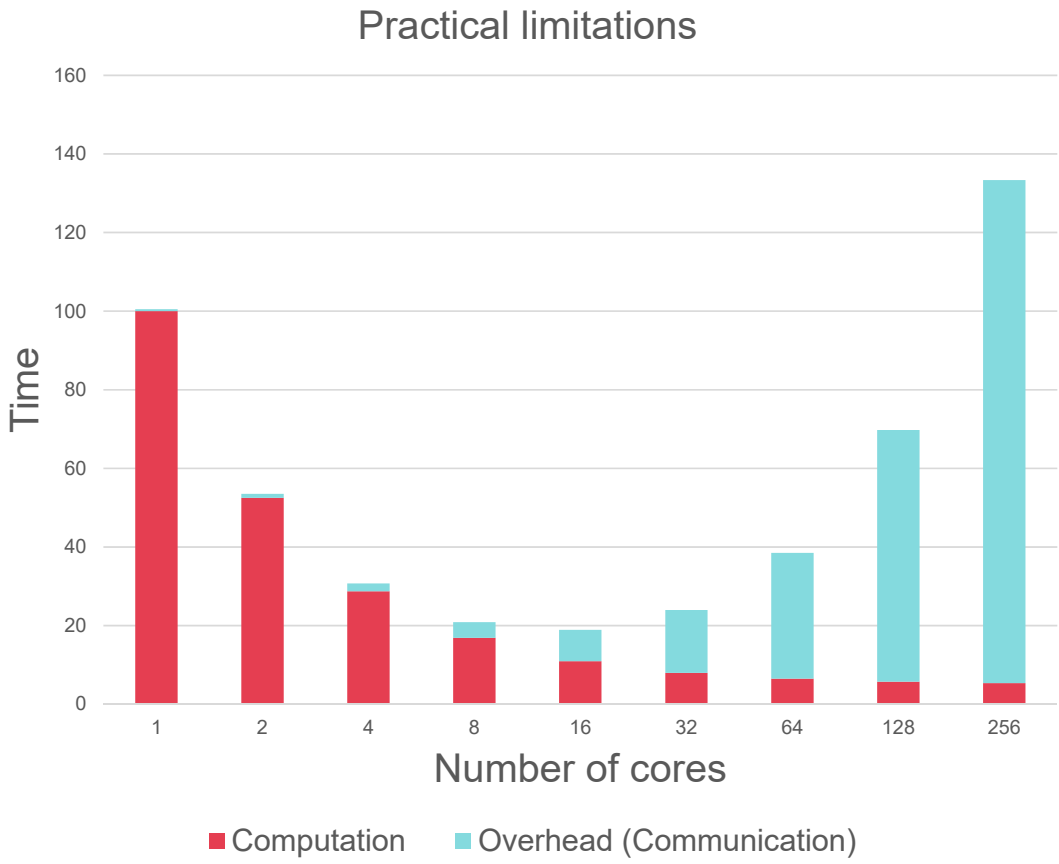

# What do you need to know in order to use it?

- Basic Linux commands.
- Scheduler.
- Theoretical limits.
- Practical limits.
- **Responsible usage.**

- Do not run applications on submission nodes.
- Do not submit unnecessary jobs.
- Do not request unnecessary resources.
- Do not connect directly to compute nodes.
- Try to optimize the use of the resources.

# Almesbar - Accounts, resource allocation and file sharing

- **User name and access**
- HPC Projects
- Jobs
- Queues
- Storage
- License applications

You will be using your KU Employee ID as your user name for accessing the HPC cluster.

You will be using SSH keys to identify yourself and access the HPC cluster.

SSH keys allow secure, password-less access to remote systems. They are very easy to set up and provide a greater level of security than using a password, while being convenient too.

We will discuss how to create and use your keys in the coming sections.

# Almesbar - Accounts, resource allocation and file sharing

- User name and access
- **HPC Projects**
- Jobs
- Queues
- Storage
- License applications

Resource allocation and file sharing on Almesbar HPC cluster is based on HPC projects.

Every single user account will be associated with at least one HPC project.

An HPC project will also be a shared space where project members can share files and submit their jobs.

Ultimately, every job has to be accounted against an HPC project.

Fairshare is applied to the HPC projects at first level.

# Almesbar - Accounts, resource allocation and file sharing

- User name and access
- HPC Projects
- **Jobs**
- Queues
- Storage
- License applications

Jobs number:

- Limit of 1000 jobs per user on **prod** and **gpu** queues
- Limit of 1 job per user on **devel** queue

Jobs size:

- Wall time limitations on each queue.
- There is no limit on the number of cores used by a job.

# Almesbar - Accounts, resource allocation and file sharing

- User name and access
- HPC Projects
- Jobs
- **Queues**
- Storage
- License applications

Access is restricted for train queue.

Research computing department uses train queue for compiling, installing, maintaining applications as well as for troubleshooting with the HPC users.

# Almesbar - Accounts, resource allocation and file sharing

- User name and access
- HPC Projects
- Jobs
- Queues
- **Storage**
- License applications

There will be quotas applied to the different storage places.

By default:

- Home directory (`/home/kunet.ae/username`): 50 GB
- Project directory (`/l/proj/projectname`): 2TB

# Almesbar - Accounts, resource allocation and file sharing

- User name and access
- HPC Projects
- Jobs
- Queues
- Storage
- **License applications**

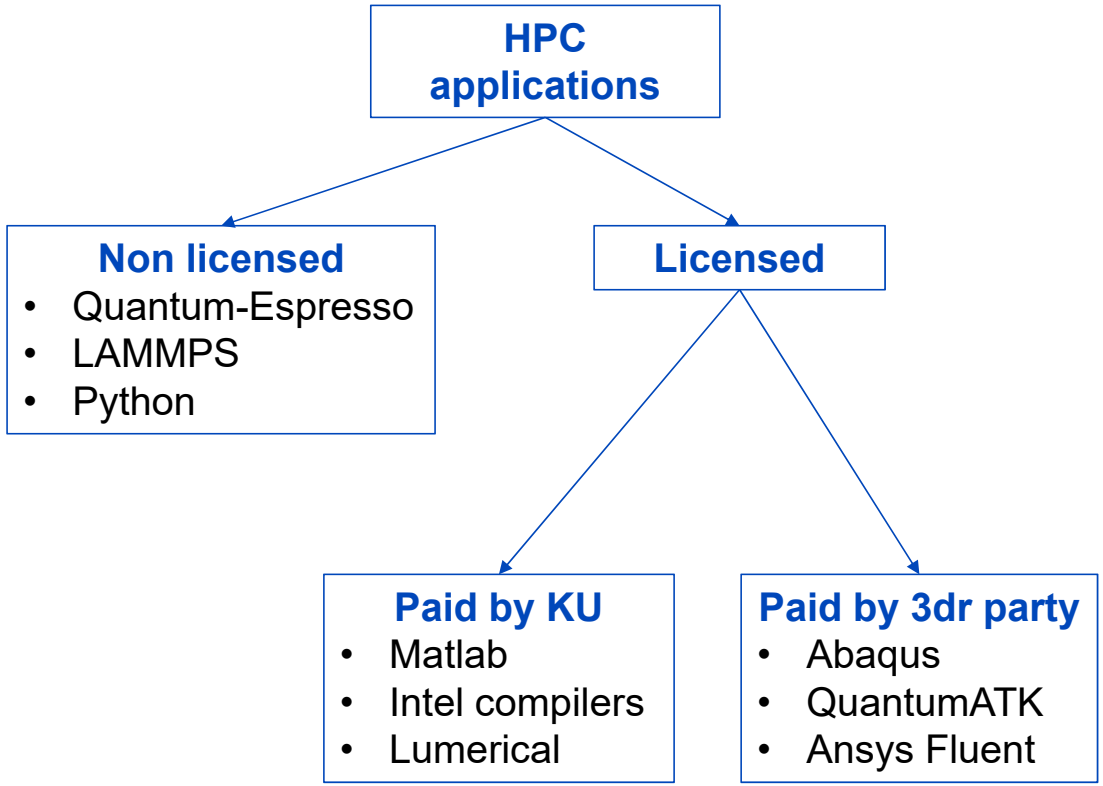

# Almesbar configuration

## Access configuration

- Queues (partitions)
- Modular environment
- Scheduler syntax

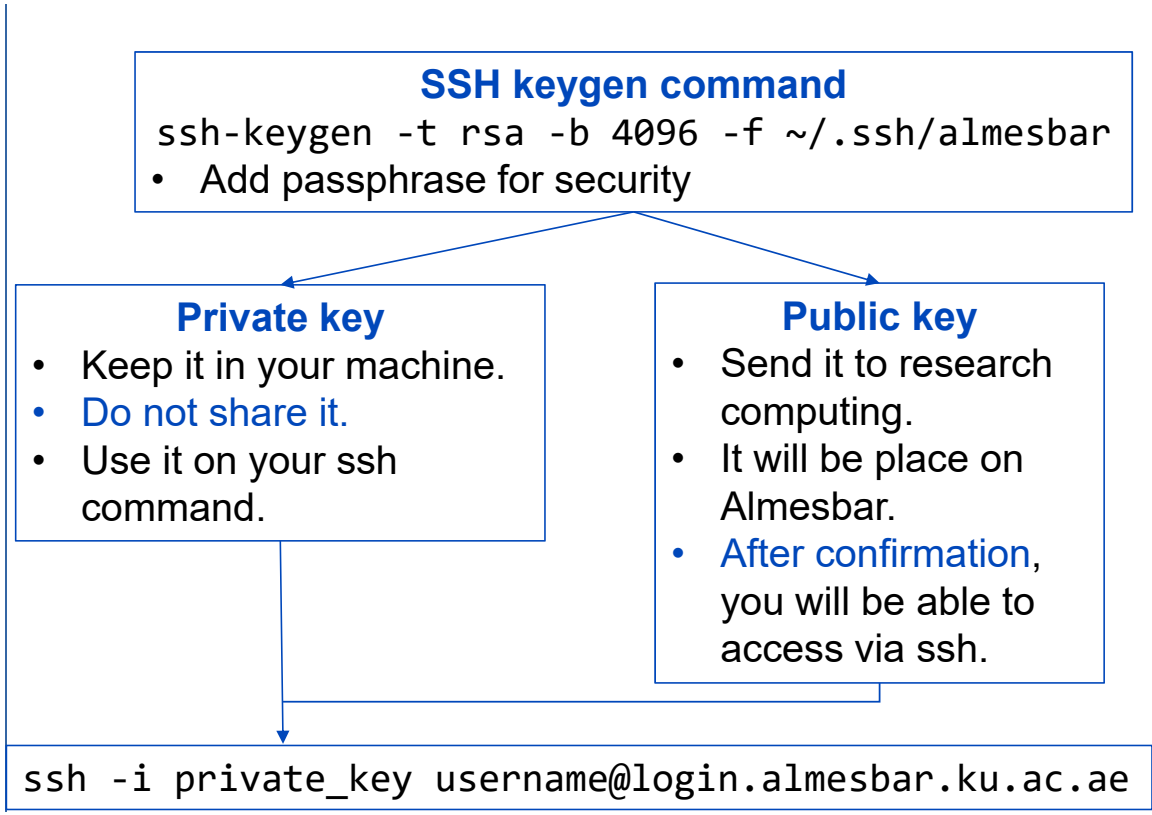

# Almesbar configuration

- Access configuration
- **Queues (partitions)**
- Modular environment
- Scheduler syntax

| Production | |
|---|---|
| Name | prod |
| Purpose | Production jobs using CPU cores. |
| Number of nodes | 80 |
| Limits | Default wall time: 2h<br>Maximum wall time: 48h<br>Max number of jobs per user: 1000 |
| Access | Open |

# Almesbar configuration

- Access configuration
- **Queues (partitions)**
- Modular environment
- Scheduler syntax

| Development | |
|---|---|
| Name | devel |
| Purpose | For interactive and testing purposes, i.e, compiling, configuring and testing applications. Running GUI applications. |
| Number of nodes | 4 |
| Limits | Default wall time: 1h<br>Maximum wall time: 8h<br>Max number of jobs per user: 1 |
| Access | Open |

# Almesbar configuration

- Access configuration
- **Queues (partitions)**
- Modular environment
- Scheduler syntax

| GPU | |
|---|---|
| Name | gpu |
| Purpose | Production jobs using GPU cores. |
| Number of nodes | 3 |
| Number of GPUs per node | 4 |
| Limits | Default wall time: 2h<br>Maximum wall time: 48h<br>Max number of jobs per user: 1000 |
| Access | Open |

# Almesbar configuration

- Access configuration
- **Queues (partitions)**
- Modular environment
- Scheduler syntax

| Training | |
|---|---|
| Name | train |
| Purpose | Used by Research Computing department for compiling, installing, maintaining applications as well as for troubleshooting with the HPC users. |
| Number of nodes | 4 |
| Limits | No |
| Access | Restricted |

# Almesbar configuration

- Access configuration
- Queues (partitions)
- **Modular environment**
- Scheduler syntax

On a complex computer system, on which it is necessary to make available a wide choice of software packages in multiple versions, it can be quite difficult to set up the user environment so as to always find the required executables and libraries.

This is particularly true where different implementations or versions use the same names for files.

Environment modules provide a way to selectively activate and deactivate modifications to the user environment which allow particular packages and versions to be found.

Our modular environment is hierarchical and it has three levels:

Core, Compiler and MPI

# Almesbar configuration

- Access configuration
- Queues (partitions)
- **Modular environment**
- Scheduler syntax

## Useful commands:

| Command | Example | Description |
|---|---|---|
| module avail | module avail | Lists available modules in the system.<br>(D) means default.<br>(L) means it is already loaded.<br>Press Q to exit the list at any moment. |
| module list | module list | List your currently loaded modules. |
| module load module_name | module load matlab/R2021a | Loads the module for Matlab version R2019b |
| module show module_name | module show matlab/R2021a | Displays information about one or more module files. |
| module spider string | module spider wrf | List all module that contain the "string“ with details on how to load them. |
| module swap old_module new_module | module swap python/2.7.3 python/3.6.0 | Switches similar/conflicting modules |

# Almesbar configuration

- Access configuration
- Queues (partitions)
- **Modular environment**
- Scheduler syntax

**Core level:**

- Modules that are not dependent on Compiler or MPI implementation.
- Most of the installed applications that don't require to be compiled seats in this level. Also the compilers can be found in this level.
- This is the only level that is visible for `module avail` command after logging in the system.
- Only if user loads a module that is a compiler, for example `gcc/9.3` , the next level in the hierarchy (Compiler) will be available and became visible with module avail.

# Almesbar configuration

- Access configuration
- Queues (partitions)
- **Modular environment**
- Scheduler syntax

**Core level:**

```
[username@login-3 ~]$ module avail
-------------------------- /apps/ku/modulefiles/Core --------------------------
   abaqus/2019            cuda/11.3               miniconda/3
   abaqus/2021    (D)     gcc/9.3                 molpro/2021.1.0
   ansys/2021R1           lumerical/2021R1        starccm/16.02.009
   ansysem/2021R1         matlab/R2020b
   comsol/56              matlab/R2021a    (D)
```

# Almesbar configuration

- Access configuration
- Queues (partitions)
- **Modular environment**
- Scheduler syntax

**Compiler level:**

- This level becomes available only if the user loads a compiler module, like `gcc/9.3` from the Core level.
- Then, all the applications compiled with that compiler (`gcc/9.3`) that don't require MPI to be compiled, will be listed here. This include MPI implementations like `openmpi` or `mvapich` for example.
- Only if user loads a module that is an MPI implementation, for example `openmpi/4.0` , the next level in the hierarchy (MPI) will be available and became visible with `module avail` command.

# Almesbar configuration

- Access configuration
- Queues (partitions)
- **Modular environment**
- Scheduler syntax

## Compiler level:

```
[username@login-3 ~]$ module load gcc/9.3
[username@login-3 ~]$ module avail

-------------------- /apps/ku/modulefiles/Compiler/gcc/9.3 --------------------
   hdf5/1.10.6         libpng/1.2.12     ucx/1.8.0
   jasper/1.701.0      openblas/0.3.7    wannier90/3.1.0
   libfabric/1.10.1    openmpi/4.0       zlib/1.2.11

------------------------- /apps/ku/modulefiles/Core -------------------------
   abaqus/2019           cuda/11.3              miniconda/3
   abaqus/2021    (D)    gcc/9.3          (L)   molpro/2021.1.0
   ansys/2021R1          lumerical/2021R1       starccm/16.02.009
   ansysem/2021R1        matlab/R2020b
   comsol/56             matlab/R2021a    (D)
```

# Almesbar configuration

- Access configuration
- Queues (partitions)
- **Modular environment**
- Scheduler syntax

**MPI level:**

- This level becomes available only if the user loads an MPI implementation module, like `openmpi/4.0` from the Compiler level.
- Then, all the applications and libraries compiled with that compiler/MPI pair ( `gcc/9.3` and `openmpi/4.0` will be listed here.

# Almesbar configuration

- Access configuration
- Queues (partitions)
- **Modular environment**
- Scheduler syntax

## MPI level:

```
[username@login-3 ~]$ module load openmpi/4.0
[username@login-3 ~]$ module avail

---------------- /apps/ku/modulefiles/MPI/gcc/9.3/openmpi/4.0 ----------------
   elpa/2021.05.001        netcdf/4.7.3       vasp/6.2.0-wannier90-3.1.0
   fftw/3.3.8              phdf5/1.10.6       wps/4.0
   lammps/29Oct2020        pnetcdf/1.12.1     wrf/4.0
   netcdf-cxx/4.3.1        qe/6.7.0
   netcdf-fortran/4.5.2    scalapack/2.1.0

-------------------- /apps/ku/modulefiles/Compiler/gcc/9.3 --------------------
   hdf5/1.10.6         libpng/1.2.12     ucx/1.8.0
   jasper/1.701.0      openblas/0.3.7    wannier90/3.1.0
   libfabric/1.10.1    openmpi/4.0  (L)  zlib/1.2.11

------------------------- /apps/ku/modulefiles/Core -------------------------
   abaqus/2019           cuda/11.3              miniconda/3
   abaqus/2021    (D)    gcc/9.3          (L)   molpro/2021.1.0
   ansys/2021R1          lumerical/2021R1       starccm/16.02.009
   ansysem/2021R1        matlab/R2020b
   comsol/56             matlab/R2021a    (D)
```

# Almesbar configuration

- Access configuration
- Queues (partitions)
- Modular environment
- **Scheduler syntax**

## Job submission control

| Command | Example | Description |
|---|---|---|
| salloc | salloc -n 1 | Obtain a job allocation. |
| sbatch | sbatch my_script.sh | Submit a batch script for later execution. |
| srun | srun -n 1 executable | If run from command line directly, srun will first create a resource allocation in which to run the parallel job.<br>srun can be run from within sbatch script or salloc shell |
| sinfo | sinfo -p prod | Display partition information |
| | sinfo -N | Display nodes information |
| squeue | squeue -p prod | Display information about jobs |
| scontrol | scontrol show job 3420 | Show active job details |
| sacct | sacct -j 3420 | Show accounting information about active and completed jobs |
| scancel | scancel 3420 | Send a signal to kill the job |

# Almesbar configuration

- Access configuration
- Queues (partitions)
- Modular environment
- **Scheduler syntax**

## Configuring a job

| Option | Example | Description |
|---|---|---|
| #SBATCH | | Scheduler directive to be used on job configuration scripts |
| --nodes | --nodes=2 | Will request tasks be run across 2 nodes |
| --ntasks | --ntasks=2 | Will start 2 MPI tasks |
| --ntasks-per-node | --ntasks-per-node=1 | Will start 1 task per requested node |
| --cpus-per-task | --cpus-per-task=10 | Will request 10 cores per task |
| --account | --account=project_name | The project your core hours will be 'charged' to. |
| --partition | --partition=prod | Submits job the specified partition. On Almesbar HPC cluster, the default queue is “prod" |
| --output | --output=file_name | Writes job output to “file_name”, %j can be used to include the JOBID, e.g. result.%j.out |
| --error | --output=file_name | Writes job errors to “file_name”, %j can be used to include the JOBID, e.g. result.%j.err |
| --job-name | --job-name=job_name | Assigns the specified name “job_name” to the job |
| --time | --time=DD-HH:MM:SS | Requests wall time limit |
| --mem | --mem=512MB | Requests a specific amount of memory for your job |
| --x11 | salloc -n 1 --x11<br>srun -n 1 --x11 xterm | Enables X11 graphical forwarding in interactive jobs |
| --pty | srun -n 1 --pty /bin/bash | Submits an interactive job that will create a terminal when the job starts. |
| --exclusive | --exclusive | Puts the host/s running your job into exclusive execution mode. |

# Our support

- **Pipeline optimization.**
- Programming consultancy.
- Code optimization.
- Build, install, maintain and test scientific applications.
- Training.
- Knowledge Base.
- Contact email.
- Location.

- Job dependencies

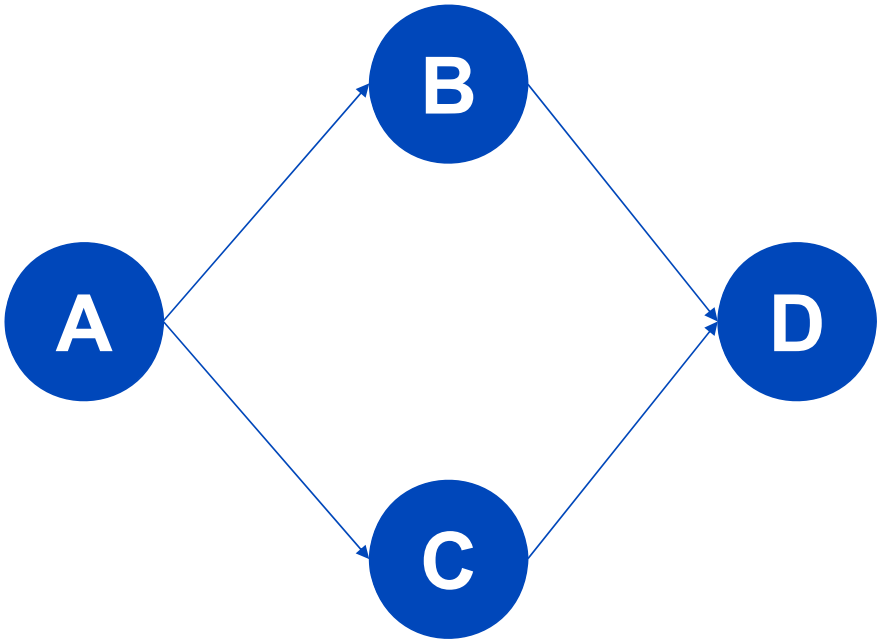

# Our support

- **Pipeline optimization.**
- Programming consultancy.
- Code optimization.
- Build, install, maintain and test scientific applications.
- Training.
- Knowledge Base.
- Contact email.
- Location.

- Array job: different parameters, but same executable.

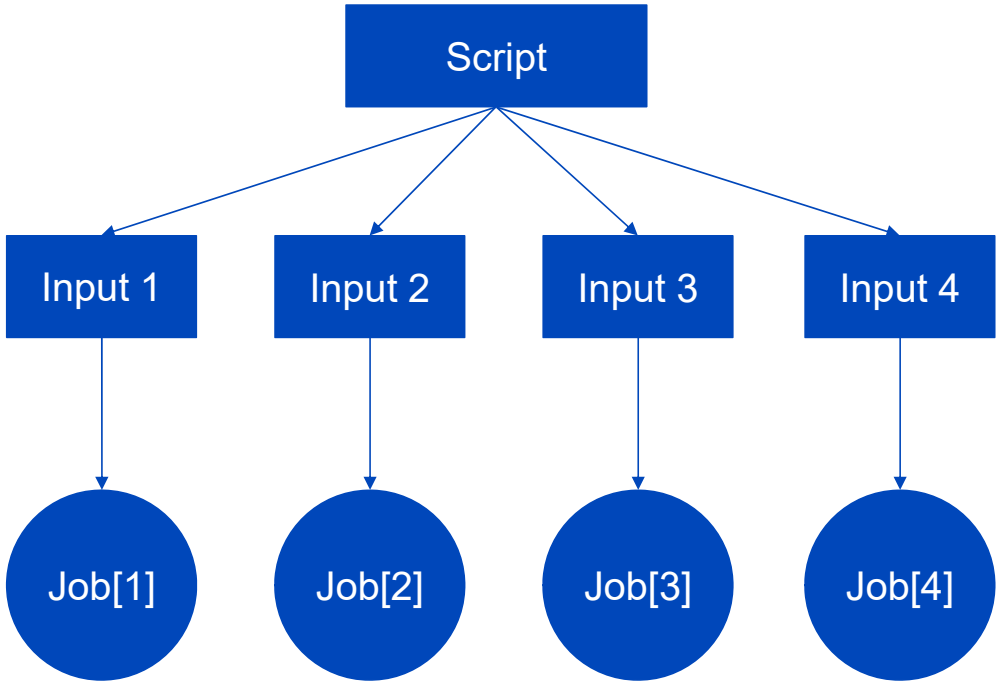

# Our support

- Pipeline optimization.
- **Programming consultancy.**
- Code optimization.
- Build, install, maintain and test scientific applications.
- Training.
- Knowledge Base.
- Contact email.
- Location.

# Our support

- Pipeline optimization.
- Programming consultancy.
- **Code optimization.**
- Build, install, maintain and test scientific applications.
- Training.
- Knowledge Base.
- Contact email.
- Location.

Code Segment 1

```
N = 10e3;

x(1)=1000;
for k=2:N
    x(k)=1.05*x(k-1);
end

Time = 0.1409 sec
```

Code Segment 2

```
N = 10e3;
x=zeros(N,1);
x(1)=1000;
for k=2:N
    x(k)=1.05*x(k-1);
end

Time = 0.00024 sec
```

# Our support

- Pipeline optimization.
- Programming consultancy.
- Code optimization.
- **Build, install, maintain and test scientific applications.**
- Training.
- Knowledge Base.
- Contact email.
- Location.

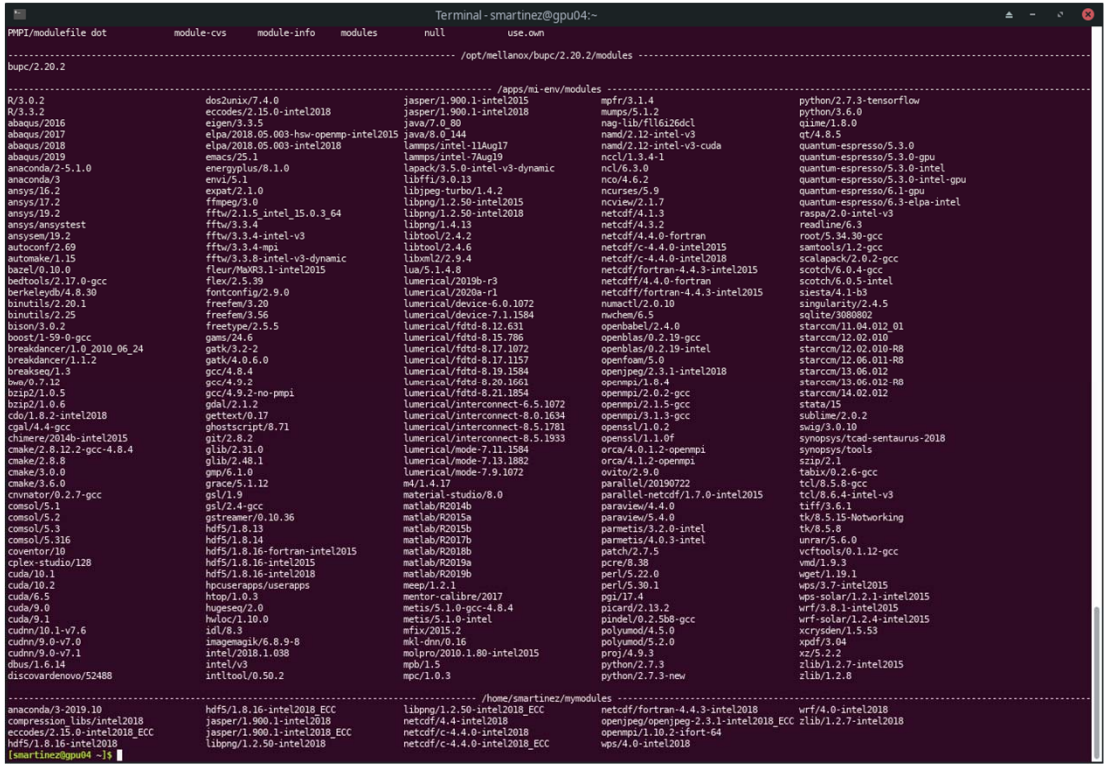

# Our support

- Pipeline optimization.
- Programming consultancy.
- Code optimization.
- Build, install, maintain and test scientific applications.
- **Training.**
- Knowledge Base.
- Contact email.
- Location.

- Research Computing **Introductory Session**.
- Introduction to Linux command line.
- **High Performance Computing** on Almesbar.
- **Numpy** - Numerical Python.
- **Matplotlib** - Plotting with Python.

# Our support

- Pipeline optimization.
- Programming consultancy.
- Code optimization.
- Build, install, maintain and test scientific applications.
- **Training.**
- Knowledge Base.
- Contact email.
- Location.

https://teamup.com/ksquu2jki56nesnr6n

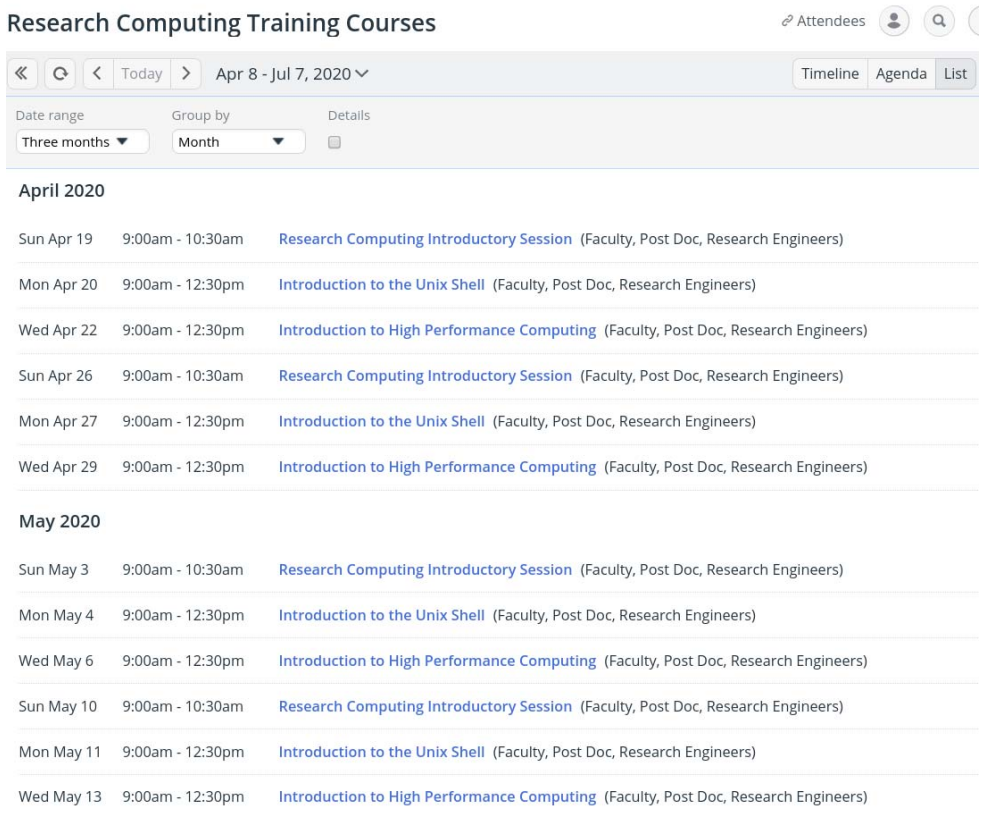

# Our support

- Pipeline optimization.
- Programming consultancy.
- Code optimization.
- Build, install, maintain and test scientific applications.
- Training.
- **Knowledge Base.**
- Contact email.
- Location.

https://wiki-researchcomputing.ku.ac.ae/

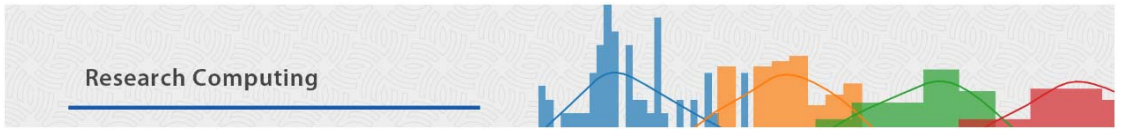


## About us

The Research Computing team provides leading edge, stable and reliable computing resources to Khalifa University researchers across all disciplines. As well as the advice, training and skills necessary to make effective use of those resources.

The mandate of the Research Computing team is to manage and support centrally provided research computing facilities at Khalifa University. These centrally provided RC facilities are the large-scale facilities such as the HPC clusters. The mandate also includes the provision of training, consultancy services and support to KU researchers on the use centrally provided resources.

The Research Computing team continually expands its services and technologies, ensuring researchers have access to a world-class computational environment.

## Services

- Support researchers by understanding their research computational needs and recommend suitable applications and other research computing resources available at the University.
- Assist researchers with grant proposals that have a strong High Performance Computing (HPC) component to describe the utilization of HPC resources in their research.
- Ensure to resolve faculty and student computing resource issues and provide programming consultancy, code optimization and resolve users' programming issues (based on Unix shell script, C/C++, Matlab, Python, MPI, OpenMP, CUDA).
- Build, install, maintain and test scientific applications on the HPC clusters.
- Develop awareness and training programs for students, faculty, and staff with regards to research computing techniques, services, and facilities, including HPC application training to help users parallelize, debug and optimize their applications.
- Conduct regular trainings workshops and tutorials, as well as one-on-one training for new users.
- Perform required regular maintenance and upgrades in coordination with external vendors.

## Quick navigation

KU HPC cluster

SAN campus HPC cluster

## Useful links

| Link | Description |
|---|---|
| Research Computing courses' calendar | Research Computing department is regularly offering a training courses on "Introduction to the Unix Shell" and "Introduction to the Unix shell". Use this link to register for future courses. |
| Software Carpentry - The Unix Shell | The content of the "Introduction to the Unix Shell" course has been taken from Software Carpentry portal. You can learn the same skills following one of their online courses. It won't take more than 4 hours. |

## Contact email

ResearchComputing@ku.ac.ae

# Our support

- Pipeline optimization.
- Programming consultancy.
- Code optimization.
- Build, install, maintain and test scientific applications.
- Training.
- Knowledge Base.
- **Contact email.**
- Location.

**ResearchComputing@ku.ac.ae**

# Our support

- Pipeline optimization.
- Programming consultancy.
- Code optimization.
- Build, install, maintain and test scientific applications.
- Training.
- Knowledge Base.
- Contact email.
- **Location.**

- **Sas Al Nakhl campus**
- **Habshan building**
- **Second floor**
- **Rooms 246 and 248**

جامعة خليفة
Khalifa University

# Thank You

ku.ac.ae